



Business Intelligence Solutions

Business intelligence (BI) solutions are all the buzz as of late, and BI developers are highly sought after. Considering the amount of data that needs to be tracked to run a business successfully, it is no wonder. When an employee has been with a company for 20 years, how will management be notified? Perhaps staffing is suffering because of vacation trends or sales need to be tracked after targeted advertising. Maybe product preordering for a sales event needs to be estimated, or who sold what and when needs to be documented for an upcoming contest.

There is no end to how much data needs to be managed, and countless hours, money, and resources are wasted in attempts to research the information, often with minimal results, multiple errors, and missed opportunities in decision making. And when more than one employee needs access to the same information, the errors are often multiplied.

With a well-designed BI solution, important data can be called up instantly in a user-friendly manner. Calculations are made with a click of a button, and reports are easily generated. No longer will that 20-year employee be unrecognized for such a long duration of loyalty and service. Staffing can be more properly managed, advertising can be better targeted to the proper demographic, and so on.

This book shows how to build a successful BI solution step-by-step. We cover the entire process from initial preparations and planning to complex layers of designing and configuring your project, and from creating reports to drafting user instructions, and releasing your project. This book is simple in its approach. If you are new to BI solutions, you will find the instructions thorough and easy to follow with clear images to demonstrate the process. Yet, it is fast-paced and rich enough in information for even the most advanced database professional to learn from.

Who Should Read This Book?

This book is for each professional who works with the many aspects of BI solutions. These include database administrators, project managers, testers, support techs, report developers, and many others.

This book is not a sales pitch for the latest features of SQL server. Nor is it focused on technologies designed only for very large companies. Instead, this book is about how small, medium, and large companies, as well as departments within those companies, can take advantage of Microsoft SQL Server's effective and inexpensive BI software. This book defines the glue that is used to bind all four of Microsoft's BI servers (MSSS, SSIS, SSAS, and SSRS) together into a BI solution.

After reading this book and working through the recommended exercises, you will have the tools to build your own BI solutions, as well as interact with other BI team members with a greater understanding of their roles within the BI solution process.

What Is a Business Intelligence Solution?

A BI solution is a collection of objects that allows data to be turned into useful information. These objects must be designed, created, tested, and ultimately approved to create a working BI solution.

When creating a BI solution, it is important first to understand what that solution consists of, how each component is combined to create the whole, and finally, how to recognize when you have achieved your goal.

Knowing where to begin is vital to the success of your project. In Figure 1-1 we have outlined eight steps to use as a guideline. We progress through each of these steps and explain them in detail throughout this book. We also develop working BI solutions in the exercises within each chapter to gain the skills necessary to complete increasingly complex solutions in your future. Chapter 2 provides an overview of the entire process.

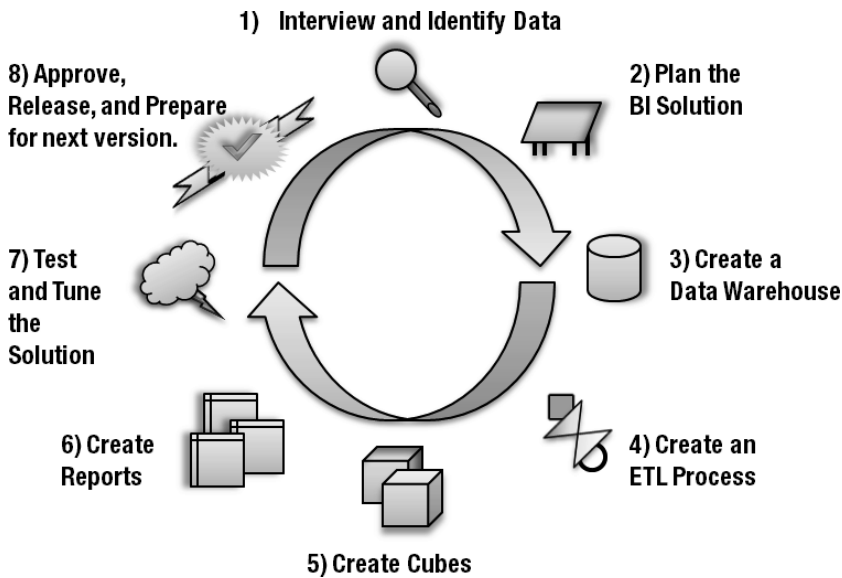


Figure 1-1. *The BI solution life cycle*

We chose to represent the tasks in Figure 1-1 as a circle, because the nature of a BI solution is one of continual change. As time goes by, a company's requirements change, the data that is available changes, and the technology to bring these two aspects together changes. Because of this, the process of creating a BI solution can often begin with the continuation of a prior solution, with each successive iteration refining and extending the current solution.

Perhaps the first step is to define the questions that your BI solution will answer. An example might be, how are our products selling? Another question might be, how often do people use our website?

One common misconception about BI solutions is that they are useful only to large corporations. This is simply not true. Clients as seemingly dissimilar as a dentist and a horse breeder will find they need to keep detailed records of important information, from patient visits to horse lineage. This information is used to determine their future plans or review past activities. Every business, group, and individual who needs to keep track of data will have questions they would like to have answered that a BI solution can provide. Formulating these questions and determining what to do with them lead us to the first step in developing a BI solution.

Step 1: Interview and Identify Data

The process of designing your solution begins with interviewing your client to determine what type of information is needed. Chapter 3 discusses the types of questions to ask and what the interview process entails.

The answers to these questions allow you to better locate the data necessary for your solution. Data can be found in many forms, and you may use one or more types to fill your requirements.

Some common data sources include the following:

- Spreadsheets
- Existing databases
- Simple text files
- Log files
- XML text files
- Paper documents

Once the data is located, the next step is to decide how much of it is relevant to your needs. You also need to decide whether your data's current location is sufficient for your BI solution's needs or whether you must copy some or all of the data to a more appropriate location. This leads us to step 2.

Step 2: Plan the BI Solution

Few developers relish creating extensive documentation before building a project. And yet, just as it is necessary for blueprints to be drawn up and approved before a home is built, projects must be planned and documented before creating a working BI solution.

In Chapter 4 we discuss creating a description of what your solution will accomplish, documenting the source and the destination objects, and beginning the formal documentation. A solution's formal document can be laid out with common tools such as Microsoft Excel or even Microsoft Word. These Excel or Word documents can then be taken back to the client for approval. Once approved, these documents will become an outline that can be worked with much like a blueprint. You then create Visual Studio projects that become the building blocks of your BI solution from these blueprints.

Step 3: Create a Data Warehouse

Your BI solution data will typically end up stored in a data warehouse database. Microsoft's SQL Server 2012 makes this very easy and cost efficient. Microsoft's SQL Server takes time and effort to master, yet the vast majority of tasks required to build your solution are performed using tools that are as simple to use as Microsoft's user-friendly Access database application.

In Chapters 4 and 5, we show how to design and implement a data warehouse database yourself, regardless of your level of experience with Microsoft's SQL Server. Various design options are demonstrated in these chapters, such as star versus snowflake dimensions and how to create fact and dimension tables. Once complete, you will understand the design differences between online transaction processing (OLTP) and data warehouse tables similar to those shown in Figure 1-2.

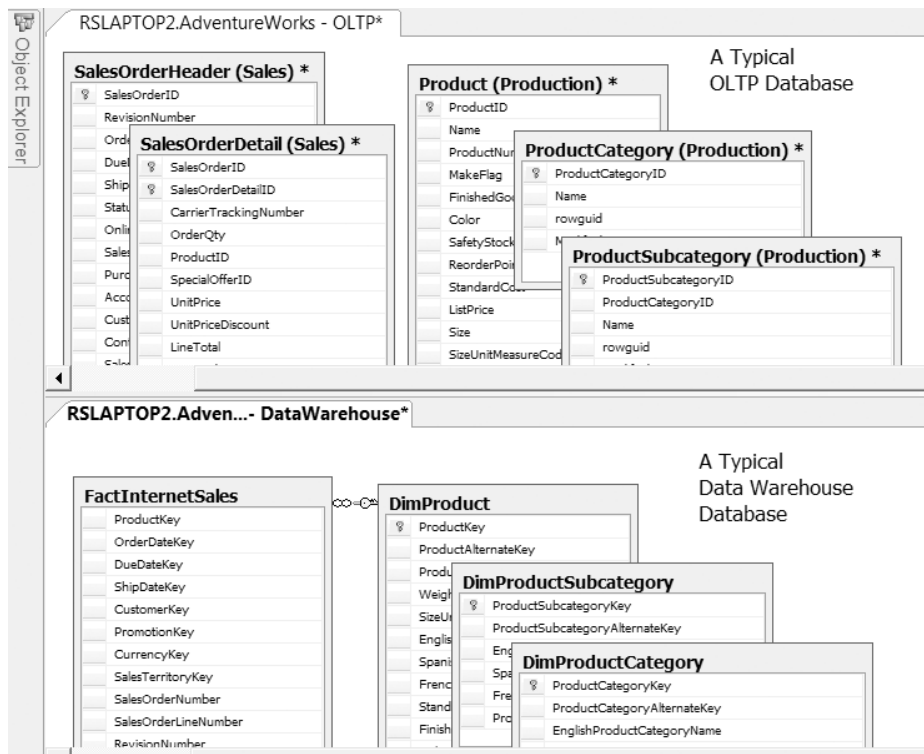


Figure 1-2. OLTP and data warehouse databases

Step 4: Create an ETL Process

Getting data from the original source to your data warehouse entails extracting the data from its original location, transforming the data to be consistent with your new data warehouse design, and loading the data into the new data warehouse location. This ETL process is discussed in great detail in Chapters 6, 7, and 8.

Although this process can be one of the most in-depth and complicated tasks in developing your BI solution, Microsoft SQL Server 2012 provides invaluable tools to help you accomplish it, saving time and simplifying the process for you. Using a combination of SQL programming and SQL Server's Integration Server (SSIS), you will create an ETL process much like the one shown in Figure 1-3.

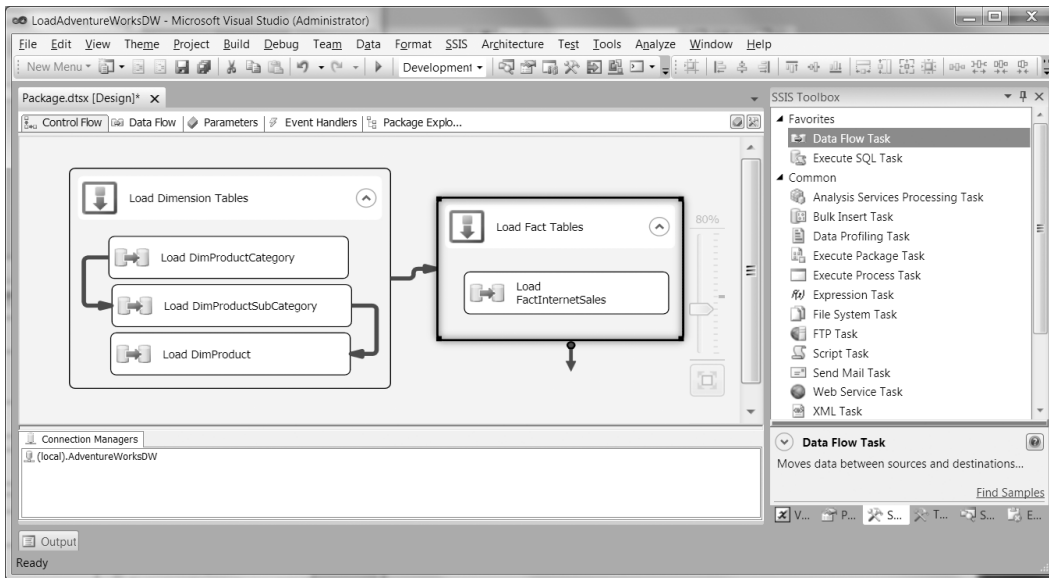


Figure 1-3. Working with SSIS

Step 5: Create Cubes

Microsoft SQL Server 2012 includes an additional high-performance server for hosting OLAP cube databases called SQL Server Analysis Services (SSAS).

Both the standard, relational data warehouse, and the SSAS cube databases have their place in BI solutions. The relational data warehouse contains a set of one or more tables and is by far the most commonly used database type. We work with this relational type of database extensively in Chapters 4 and 5. The second type of database contains one or more cubes instead of tables. You can think of these cubes as a set of report tables combined into a single object. Figure 1-4 illustrates how a cube is configured using an SSAS project in Visual Studio 2010. We discuss constructing and configuring cubes in Chapters 9 through 12.

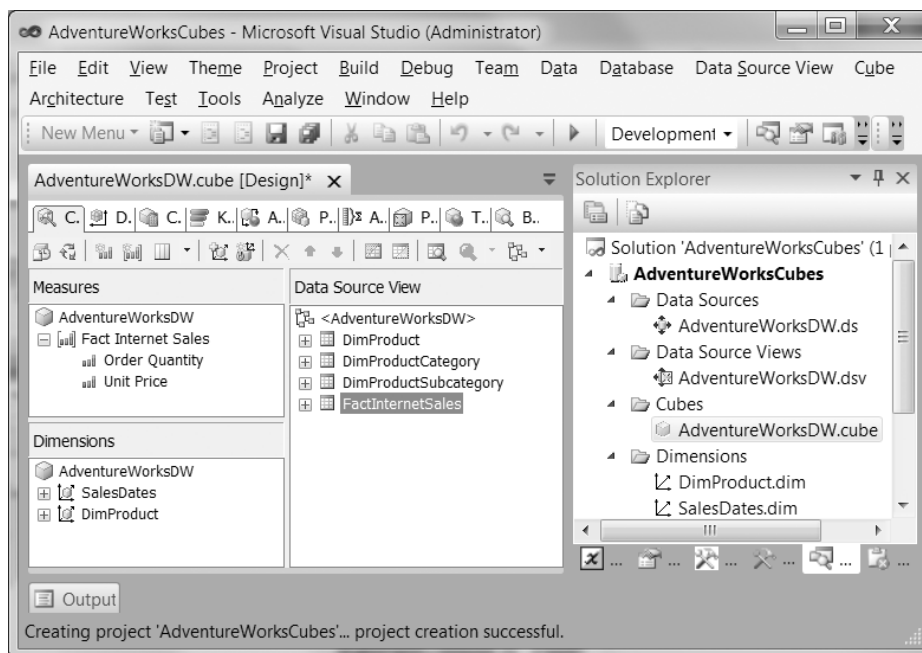


Figure 1-4. Configuring a cube in SSAS

Step 6: Create Reports

Once you have your data loaded into a data warehouse and/or cube, you need to create preliminary reports to continue your work. These may be your first reports for your BI solution, but they certainly will not be the last. The end goal of a BI solution is to convert data into usable information, and that information is routinely represented within reports.

The term *BI solution* is not very self-explanatory. It might be better if the industry as a whole changed the term *business intelligence solutions* to *business reporting solutions*. Even *make life easier on managers solutions* might be more descriptive than *business intelligence solutions*.

■ **Note** About a year ago, Randal performed a casual experiment to see how many of his co-workers within the IT industry understood what the term *BI solution* meant. As he expected, 90% did not know. Some guesses were pretty comical. A favorite was “intelligent robots for businesses.” But many guesses were nothing more than a long string of verbs in search of a definition. As you might imagine, only about 10% of his co-workers had a problem figuring out what a reporting solution was.

No matter what you call your BI solution, the most common output is a set of reports that present meaningful information to your users. You have many reporting tool options from which to choose. In this book, we focus on using the most readily available Microsoft technologies to create your BI reports, including Excel and SQL Server Reporting Services (SSRS).

Deciding what type of data source the reports will use is an important aspect of reporting. A typical pattern in the industry begins with simple solutions and moves progressively toward more complex ones over time (outlined in Figure 1-5).

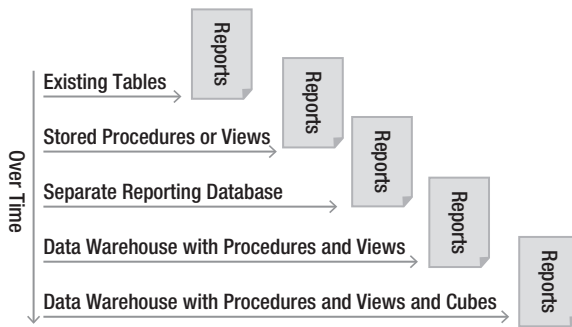


Figure 1-5. An example of how reporting data sources change over time

Many companies begin by selecting report data directly from OLTP relational tables. Quite often, they come to regret this choice when performance issues occur and maintenance costs rise. It has long been considered a poor choice to do so, yet this is still happening in businesses today.

An improvement on this design, and what is considered to be “best practice,” is to create views or stored procedures that select data from one or more OLTP tables and use these as the source for all of your reports. Many reports can then be created against a single view or stored procedure, which makes maintaining your reports much easier over time. For example, consider a scenario where a decision has been made that all tables must be renamed to start with the letters *tbl_*. All that you need to do to keep your reports working properly is change the table names in the select statements within the view or procedure to reflect the new table names, while maintaining the same output from the view or procedure. With this simple step, your reports will continue to work as they always have. Chapter 13 of this book shows how easy it is to create both views and stored procedures.

Stored procedures and views can access data in the same database, across databases, and even across different database servers. You will gain better performance, however, when you query data from a dedicated reporting database, otherwise known as a *data warehouse*. These report databases are designed to provide simple and efficient reporting. Once the data warehouse has been created, you need an ETL process to copy the data from its original locations to the new reporting data warehouse database.

■ **Note** The term *data warehouse* can have a number of meanings. In this book, a database designed for reporting with one or more centralized fact tables containing measured data such as sales quantities, with zero or more supporting dimension tables containing additional measured data descriptions, is considered a data warehouse. You may hear this type of database referred to as a *data mart*, *data silo*, *data factory*, and a host of other names. However, Microsoft documents refer to it as a data warehouse, so we do too.

Additional report performance is provided by using SSAS cubes. This performance increase, however, is at the cost of your solution becoming more complex. The most common complexity is that cube databases use different programming languages than relational databases. We discuss the most common of these programming languages, known as MDX, in Chapter 14.

To round out your report-building skills, we present report-building applications in Chapter 15. We work with Microsoft’s desktop-based reporting application, Excel 2010. Then, in Chapters 16 and 17, we create reports using Microsoft’s server-based reporting application, Reporting Services 2012.

Step 7: Test and Tune the Solution

Once you have built your first reports, you need to test those reports for accuracy, visual consistency, and performance. The most important of the three is accuracy. If the reports are slow or do not look professional, it is indeed cause for concern, but if your reports are inaccurate, your entire BI solution will fail! We cover a number of ways to plan and implement testing procedures in Chapter 18. We also include important performance-tuning techniques in Chapter 18 to insure your reports run quickly for your end users.

Step 8: Approve, Release and Prepare

At the end of the solution development cycle, you need to package and deploy your documents, scripts, databases, and reports. You also need to create user documentation, as well as train your users to use your newly developed BI solution. These topics are discussed in the last chapter of this book, Chapter 19.

Practice Exercises and More

Rather than just talking about all of these subjects, the chapters in this book offer detailed instructions on how to perform your BI solution tasks with step-by-step practice exercises that build upon each other from one chapter to the next. We created simple, easy-to-follow examples that outline key principles applicable to both large and small BI solutions.

We also offer “Learn by Doing” activities at the end of each chapter. These activities provide an outline and hints indicating which course of action to take, but they allow you a chance to practice your skills without such detailed instructions. Table 1-1 describes the exercises within this book.

Table 1-1. *Exercises in This Book*

Exercise Type	Description	Instructions
Exercises	Detailed, progressive, step-by-step instructions that correspond with the subject matter within each chapter. A complete and functioning BI solution is created by the end of this book.	Detailed instructions are included within each chapter.
Learn by Doing	A simple outline of the steps required to implement a BI task that corresponds to the subject of each chapter.	Outlined instructions are within folders included in the downloadable book content. See this book's catalog page at www.apress.com/9781430234883 .

You are given the opportunity to accomplish multiple BI tasks by the end of each chapter. The goal is to help you master the steps involved in building your own real-world BI solutions.

Downloadable Content

All example projects, exercises, and scripts have been organized into folders by chapter and compressed into zip files. This downloadable content includes all of the BI solution files and information pertaining to the locations of the original databases to make these files work.

You may at times need a hint on how to complete a task. Not to worry, help is available in the form of completed and commented solutions to each standard exercise and “Learn by Doing” exercise.

All of this and more can be found on the Apress website: www.apress.com. See the catalog page for this book at www.apress.com/9781430234883.

In addition, there is even more content available on each of the author's websites: <http://NorthwestTech.org/ProBISolutions> and www.keystrokepublications.com. Here you will find things that just could not fit within this one book such as articles, demos, templates, and videos!

Our Example Scenarios

We work on two BI solution scenarios in this book. Each scenario is based on a sample database created by Microsoft for demonstration purposes. The databases are as follows:

The Publications BI solution: The Pubs database was created in the 1980s for both Sybase's and Microsoft's SQL Server demonstrations. Pubs has a number of flaws in its design, naming conventions, and datatyping. This provides an opportunity to remedy the flaws during the creation of the data warehouse and the ETL process, just as you would find in a real-world scenario. This database also has a number of archetypal data structures useful for highlighting advanced dimensional structures. Another advantage to the Pubs database is that it is the simplest Microsoft demonstration database available. Because of all of these features, we use it as the focal point for the in-chapter practice exercises.

The Northwind Foods BI solution: Made in the early 1990s, the Northwind database is larger and slightly more complex than the Pubs database. It was also created for demonstrations by Microsoft and has numerous design flaws that are discussed and addressed in our data warehouse and ETL processes. This database is used to frame the "Learn by Doing" exercises for each chapter.

All of these databases are readily available and have been used as examples in hundreds of books. Because of this, you may already be familiar with these databases, and you can easily find additional information and code samples to enhance your understanding.

Setup Instructions

Although we have tried to keep the setup requirements as light as possible, there are still a number of complex tasks that need to be performed before you can get the full benefit of this book. You need the following items:

- A full install of SQL Server 2012 developer edition, with all of its supporting servers (SSIS, SSAS, and SSRS)
- The Pubs and Northwind databases
- Administrator-level access to SQL Server and its supporting servers (SSIS, SSAS, and SSRS)

We included setup instructions, files, and videos in a single folder called `_SetupFiles` that is included as part of the downloadable content from the Apress website, www.apress.com. Therefore, you have only one downloadable file to worry about. This folder is inside the same zip file as the exercises.

Of course, you have to unzip the file before you can use it. We include detailed instructions on how to copy it to the root of your C:\ drive in Chapter 2, but you can unzip the downloadable content anywhere you want until then. On a Windows 7 PC, the typical location would be the Downloads folder.

In Figure 1-6, we have unzipped the file and copied the resulting folders to the location described in Exercise 2-1.

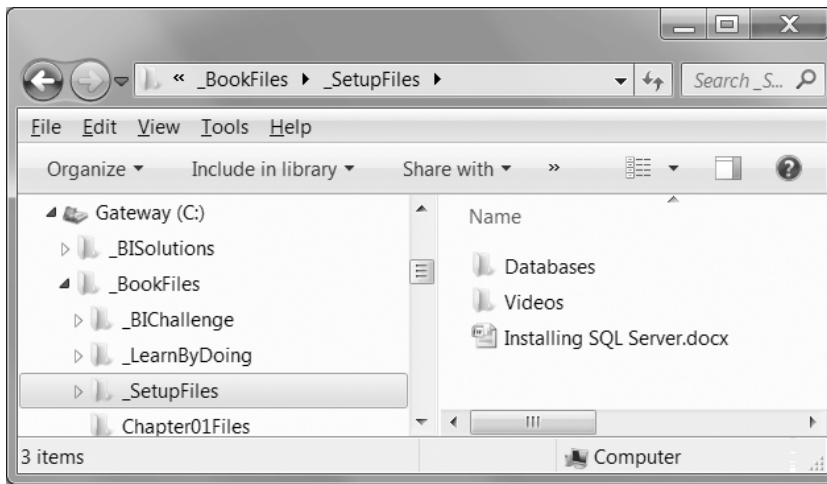


Figure 1-6. Setup files and folders

Please review the files in this folder before you start to go through this book. Full instructions are found inside the _SetupFiles folder.

■ **Tip** We have included additional videos and links that can help you tackle the installation if you still feel overwhelmed. These are found on one of the author's websites at www.NorthwestTech.org/InstallingSoftware.

Think Small, Win Big

Creating BI solutions has never been easier. The tools that many vendors offer have become more refined and user-friendly than was dreamed of a decade ago. Still, even with good and inexpensive tools, a BI solution can go horribly wrong if it is not planned and implemented properly.

In the past, a number of approaches have been attempted to ensure that BI solutions have a big impact on a business. One early approach was to include everything that was needed by the business into one master solution. These solutions often took years to complete and were not always consistent with a company's current needs by the time they were finished. This led to a number of issues that have now become widely believed misconceptions about BI solutions. These misconceptions include the following:

- They take years to implement before anything useful is available to the end users.
- They take months of planning before they even get started.
- They cost a lot of time and money.
- They are a luxury, applicable only to large companies with large budgets and large development teams.

Large and long-term solutions have their place, but they are not always necessary. Many companies can benefit immediately from small, quickly designed, and quickly developed solutions. We even go as far as to say that most BI solutions will easily fit this pattern.

A number of changes in IT over the past decade have allowed small BI solutions to become viable. The computers and the software that we run on them are more powerful and less expensive. Something as simple as a Microsoft Excel spreadsheet, for example, can now work with millions of rows at once, allowing you to create very simple BI solutions starting with that tool alone. Microsoft's SQL Server, which has always been reasonably priced, can now work with many terabytes of data, run distributed queries among a collection of servers, and comes with powerful BI tools such as Integration Services, Analysis Services and Reporting Services, at no extra cost. To see what we mean, compare earlier versions of Microsoft Excel and SQL Server. You will see that the cost to purchase these tools, without all of these new features, was roughly the same in the 1990s as it is today, not even taking into account the difference due to inflation.

The combination of more powerful computers and inexpensive software add up to a big win for small to midsize businesses. These businesses can now afford to perform BI tasks that traditionally only their larger competitors were capable of.

The following examples give an idea of how small BI reporting solutions can provide a big win to any type of business:

- Monthly sales reports for a gift shop
- Reports on a development team's projects over time
- Reports that track medication dispensed within a medical clinic
- An auto part store's inventory reports
- Reports that track support calls to a call center

Considering how reporting solutions can be beneficial to companies with 10 employees or 10,000 employees, it is no wonder that BI is such an expanding aspect of our IT industry.

Rapid Application Development for BI Solutions

Once you have established the need for BI solutions, how do you successfully plan, start, and complete them? Although there is no single answer, experience has shown that completing simple, fast, and extensible solutions are the most likely to provide the best cost-to-benefit ratio.

One of the more popular ways to initialize the development process is by using the techniques associated with rapid application development (RAD). In RAD, you start with a short planning phase, followed by a short development phase working on a simple prototype. You then test your prototype for accuracy, consistency, and performance. Once the testing phase has passed, the next step is to release the prototype for comments and prepare to start the next iteration of your solution. This next version of your solution takes comments about the existing features into account and extends the previous solution with new ones. The cycle continues, providing increasing benefit to your users over time.

RAD will not work for all projects, but it will work for a majority of them. This is one of the more successful techniques in the industry today; therefore, we focus on building solutions based on this methodology.

Moving On

In this chapter, we have outlined the steps needed to create a BI solution and discussed the subject matter covered in this book. In Chapter 2, we take a more in-depth look at the entire process by building a very simple BI solution. We start with gathering solution requirements and end with a simple, functioning prototype BI solution. It is time now to get your hands dirty and start work!

What's Next?

In each chapter, we have made our best attempt to focus on what is essential knowledge for every BI professional. We realize that this topic is much too complex for any one book and our essentials may not cover all you need to know. To help further your understanding of the topic within each chapter, we have included reading suggestions for further study.

For more information on RAD, we recommend the book *Rapid Development: Taming Wild Software Schedules* by Steve McConnell (Microsoft Press).